



Building security that thinks

Machine learning fundamentals for cybersecurity professionals


Chris Morales, Head of security analytics

chris@Vectra.ai



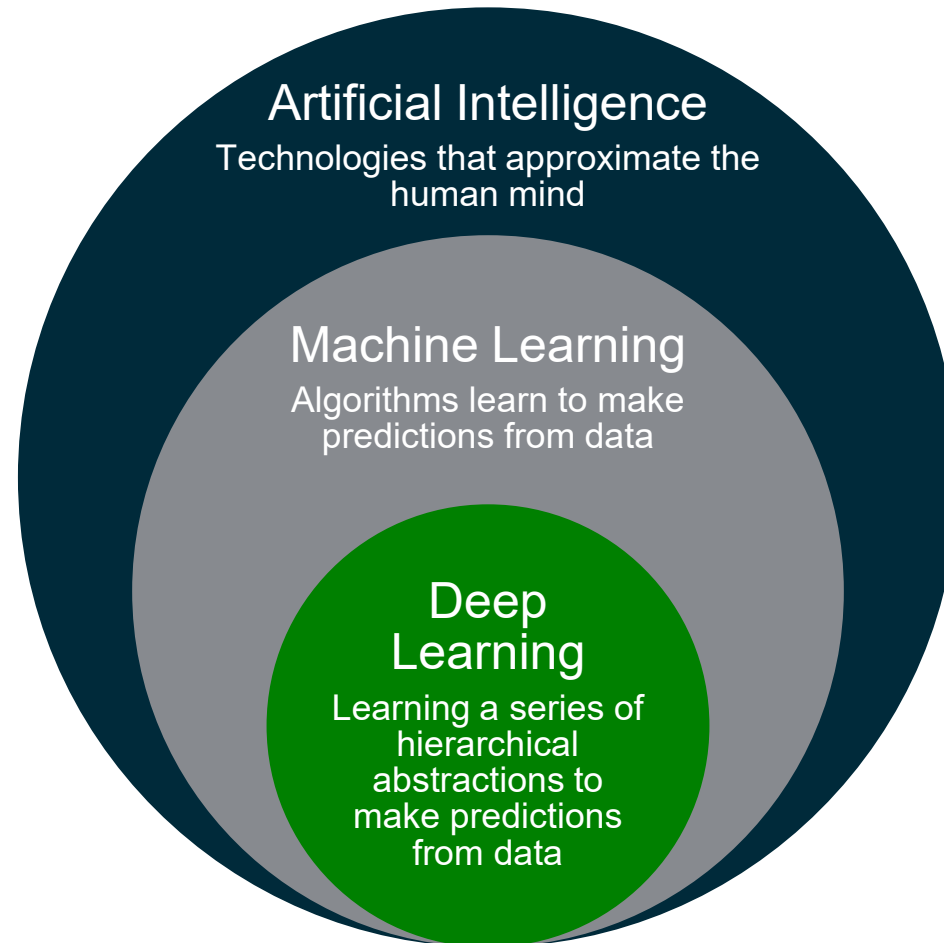
The five questions that data science answers

- Is this A or B (or C or D)?
 - Classification
- How much / How many?
 - Regression
- How is this data organized?
 - Clustering
- Is this weird?
 - Anomaly
- What action should be taken?
 - Reinforcement



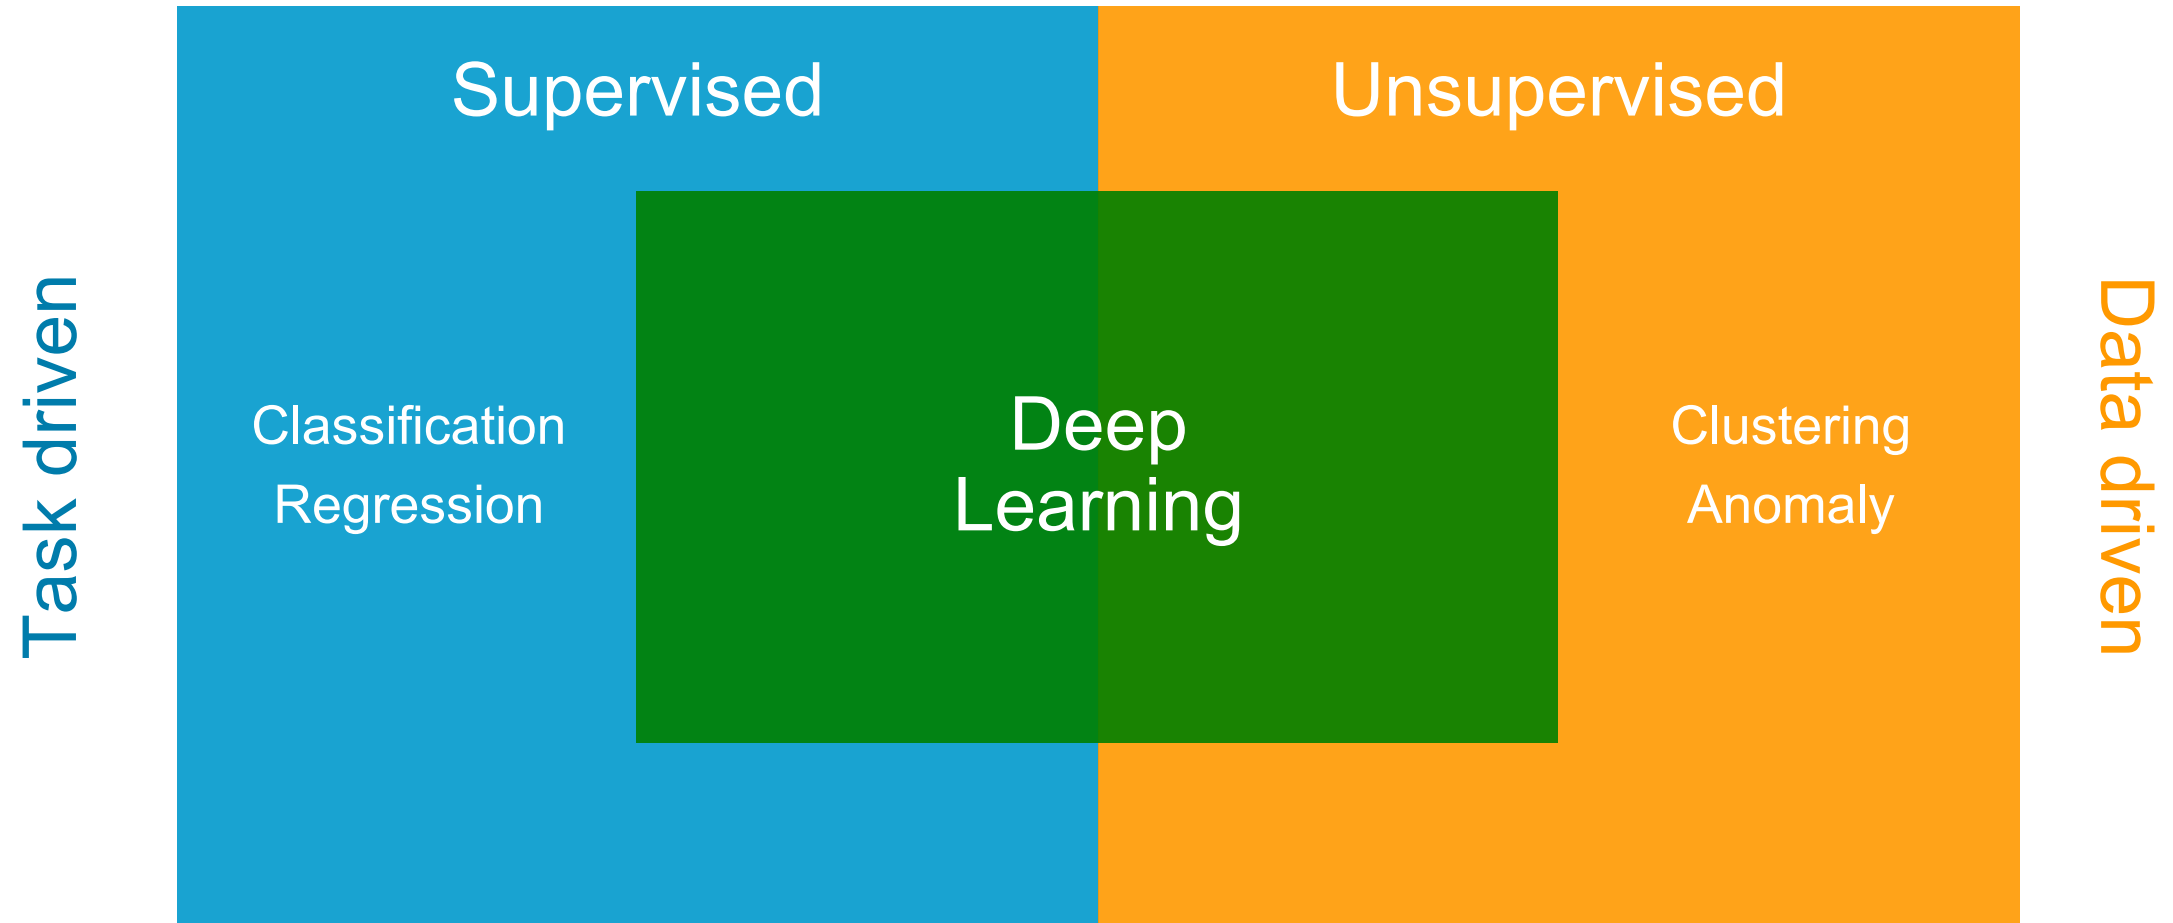


How does data science create intelligent machines?





Supervised and unsupervised machine learning






Supervised Classification - Is this A or B (or C or D)?


Predict Cat or Dog given

[Brown (T/ F) ,
Weight (0 -200 lb),
Height (0 - 3 ft)]



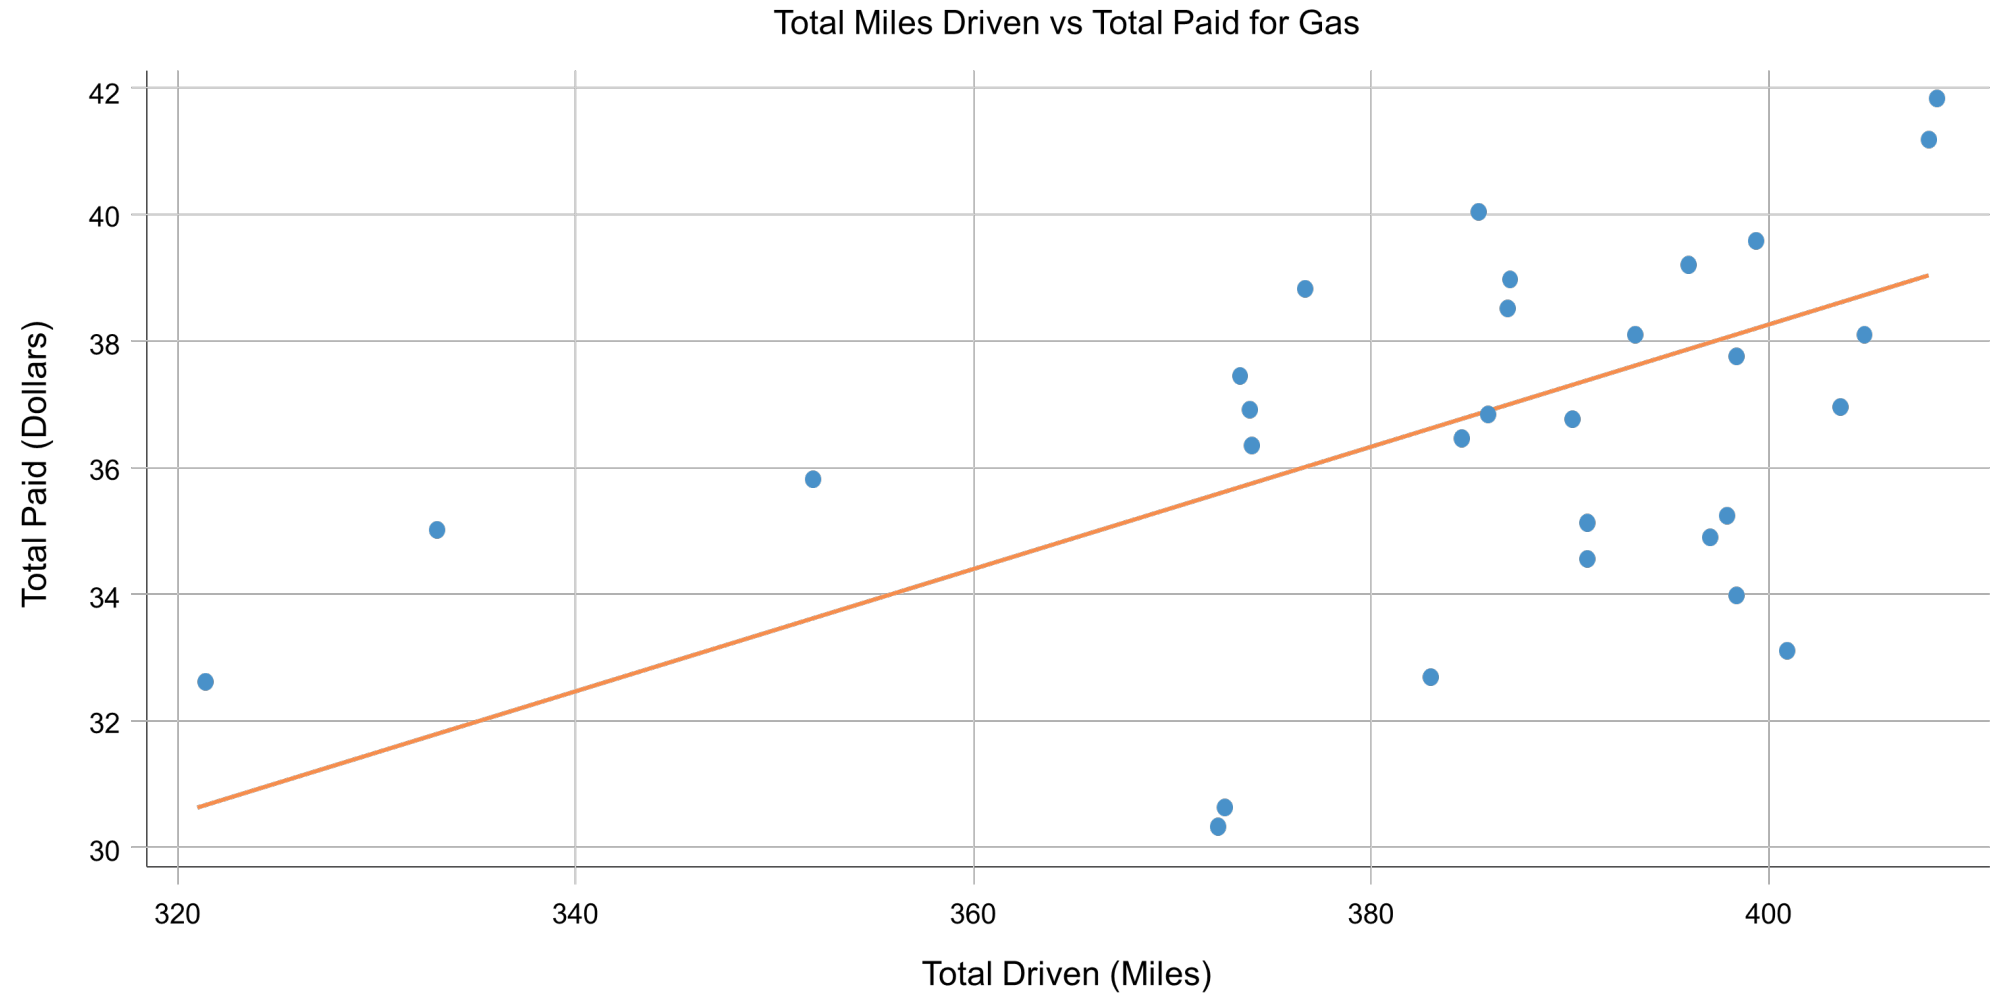


Supervised Classification - Random Forest



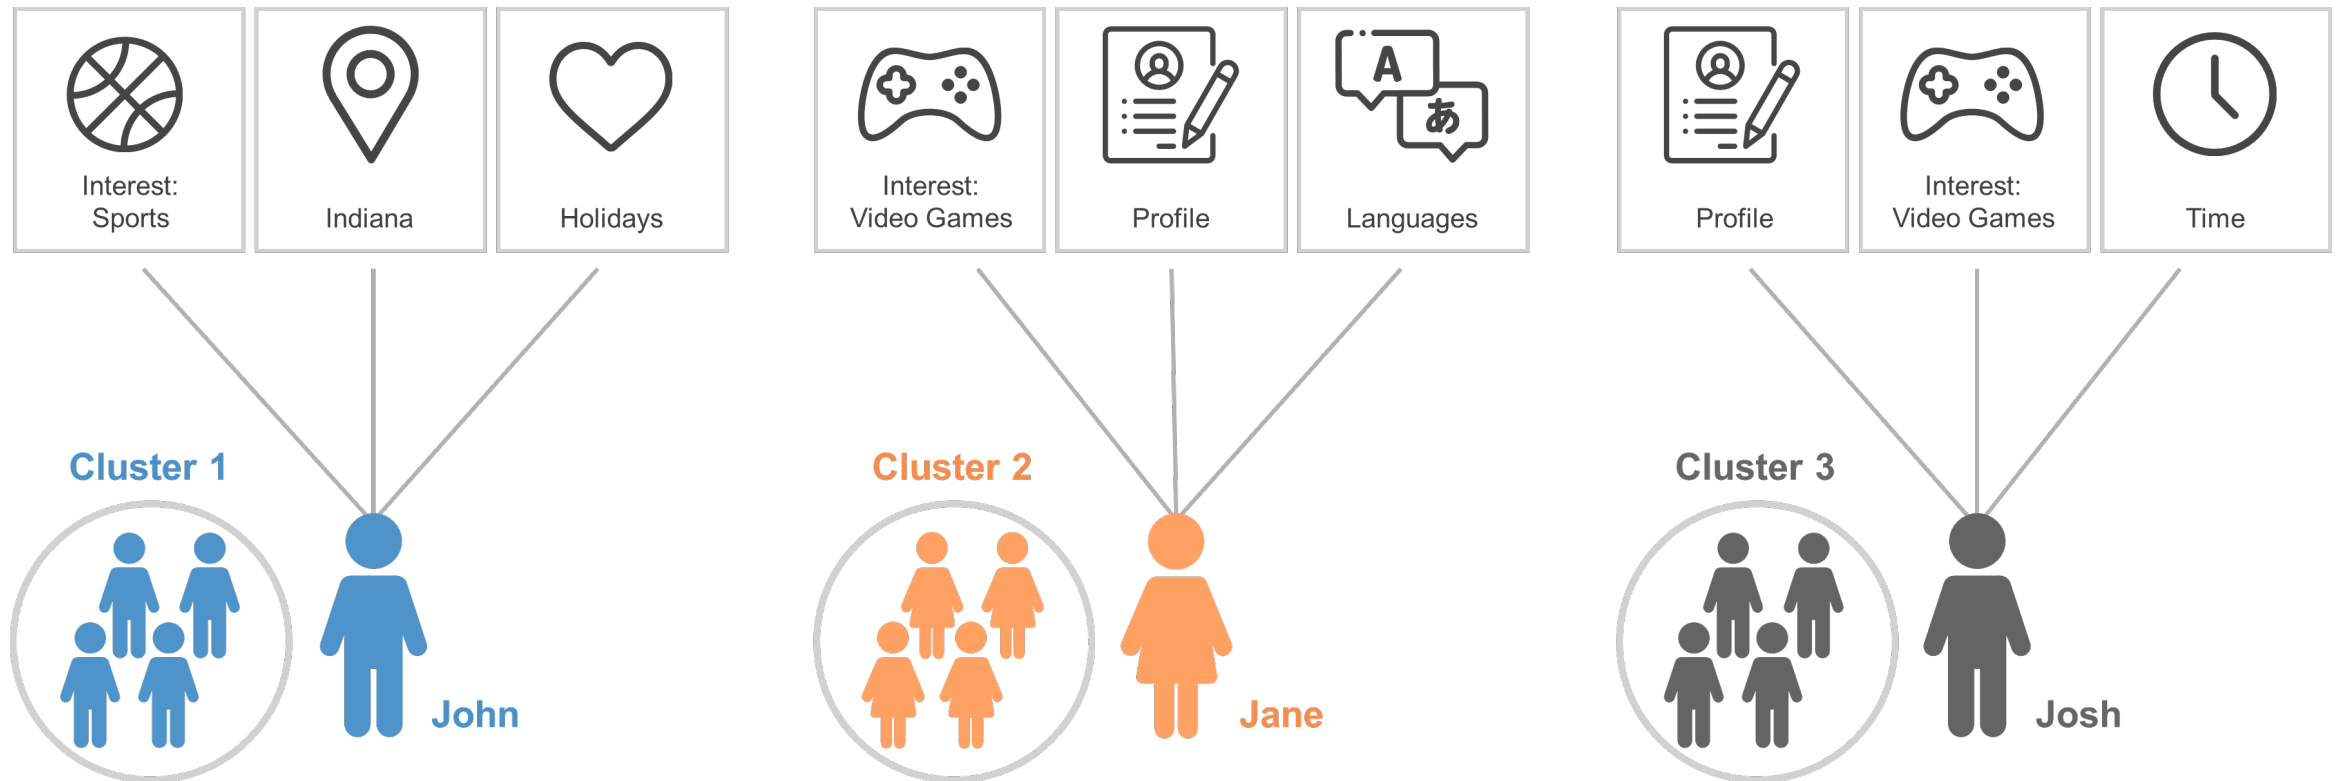


Supervised Regression - How much / How many?








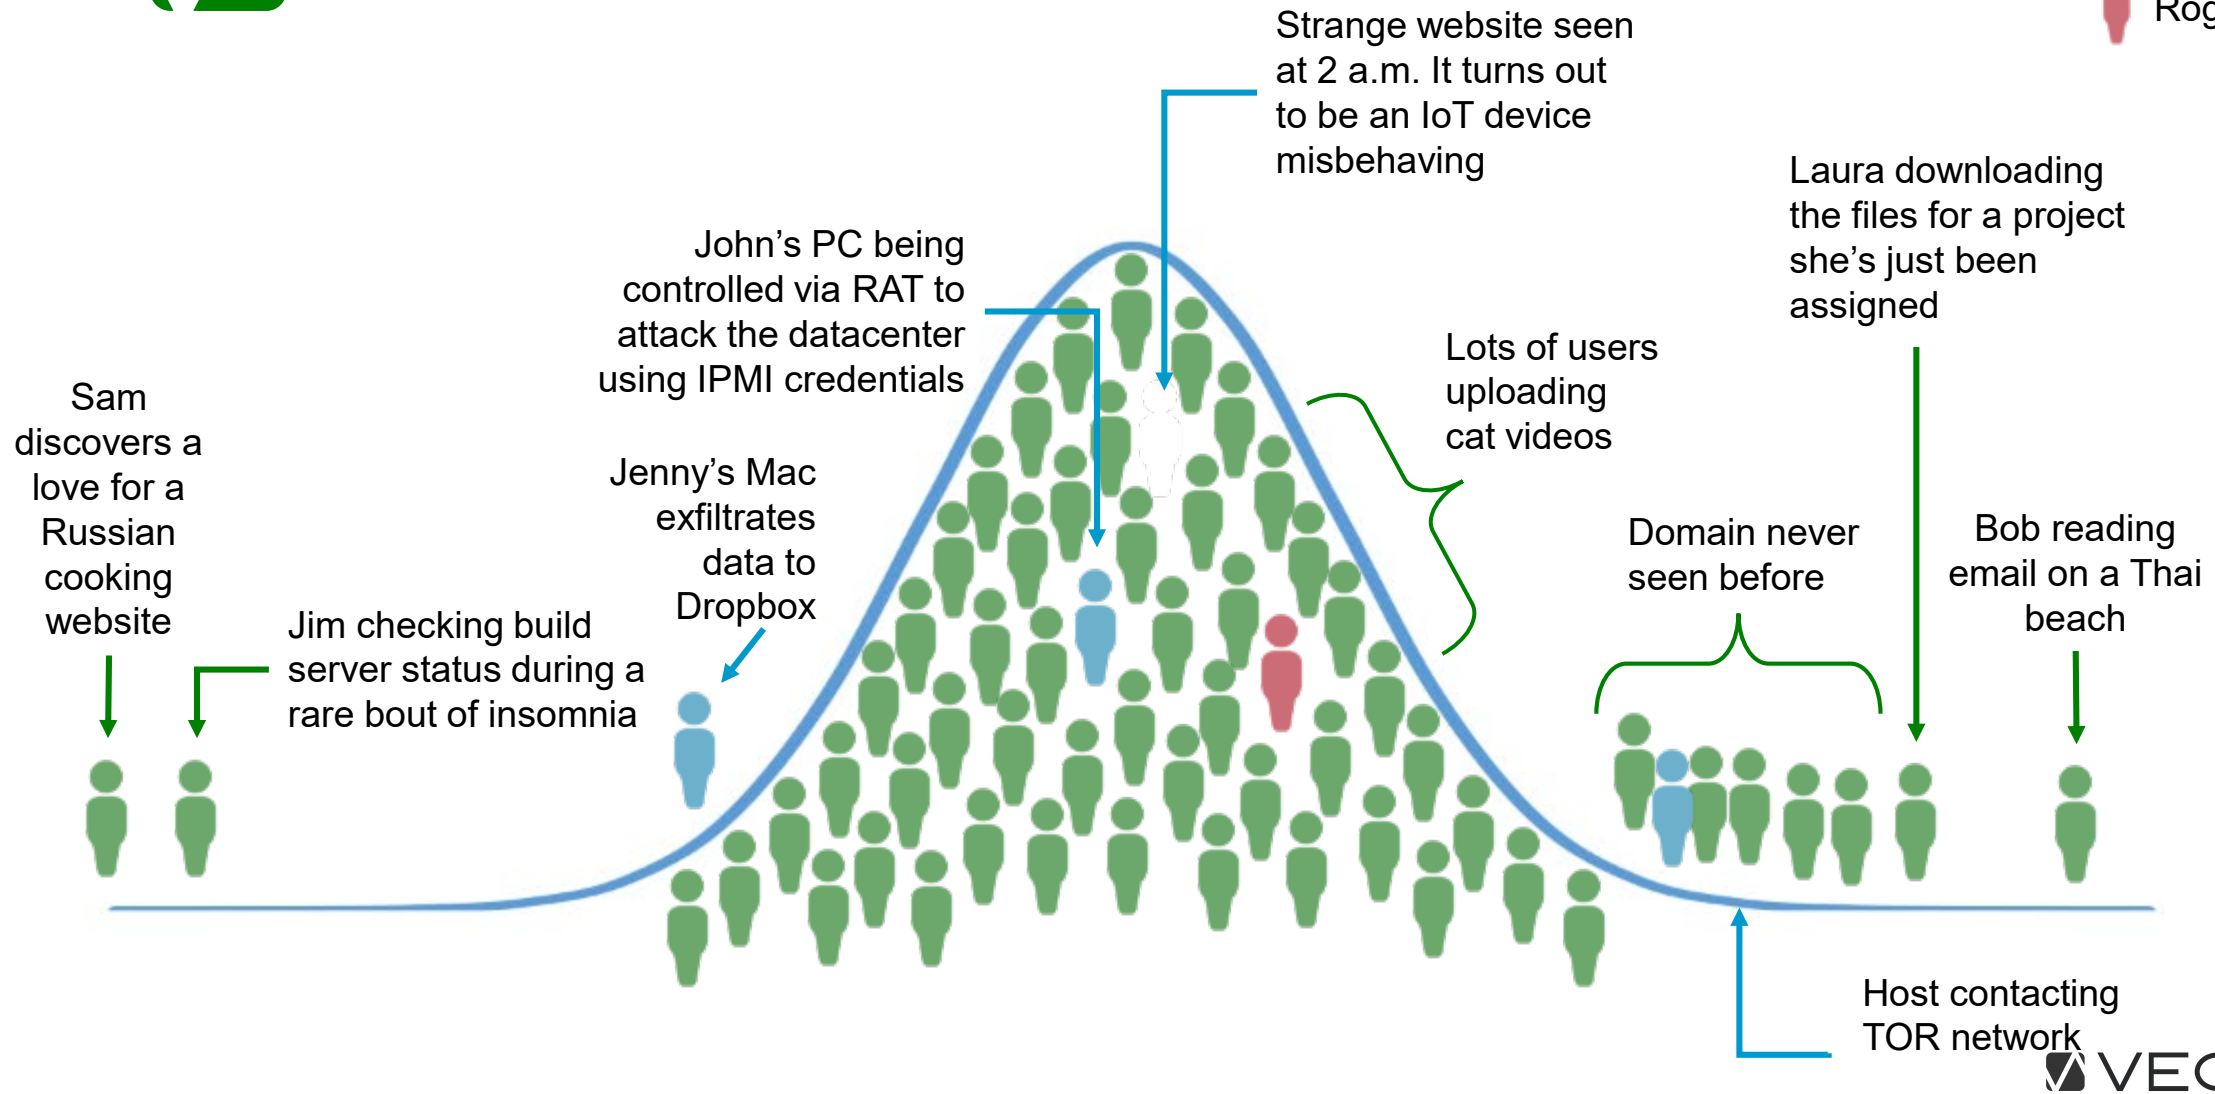
Unsupervised Clustering - How is this data organized?






Unsupervised Anomaly – Is this weird?

-  Good user
-  Attacker
-  Rogue employee




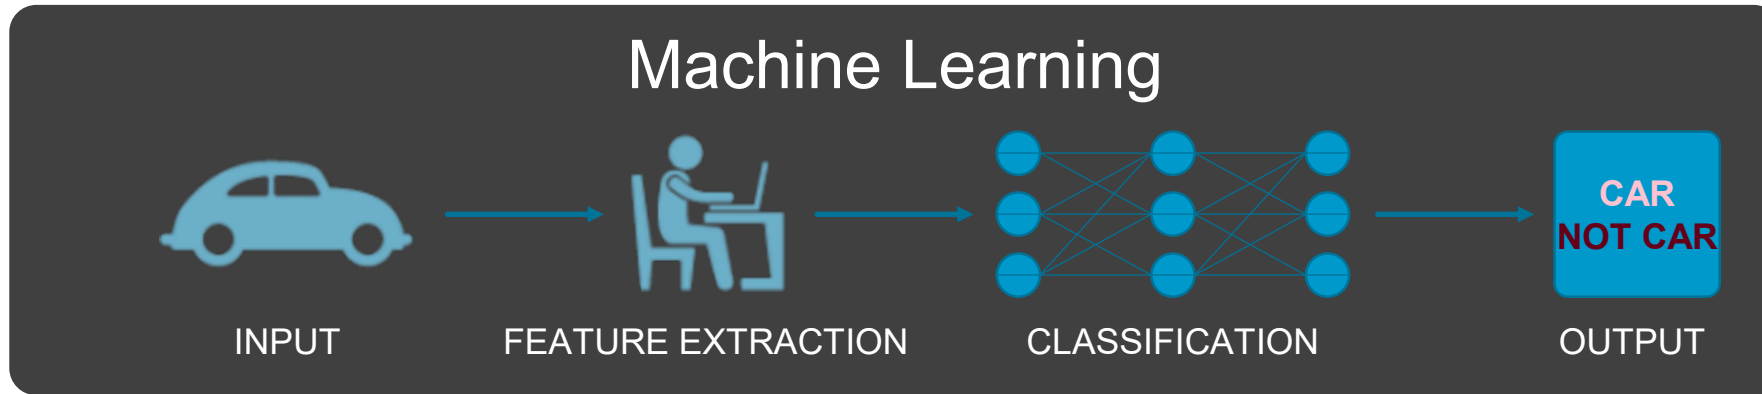


Deep learning - Using artificial neural networks





Choosing traditional machine learning or deep learning





Applying data science to threat detection

Signature



How the threat looks

Find threats that you've seen before

Snapshot in time

No local context

Data science



What the threat does


Find what all threats have in common

Learning over time

Local learning and context



Looking for what the threat does



Durable coverage

- Both novel and known attacks
- Difficult and expensive to evade

Fast, labeled coverage of known threats

- Tools
- Exploits
- Known attacker infrastructure
- Environment-specific indicators



Combine data science with security research

Attacker Behavior models

- High-fidelity detection of things attackers must do
- No signatures: find known and unknown

Security Research

- Identify, prioritize, and characterize fundamental attacker behaviors
- Validate models

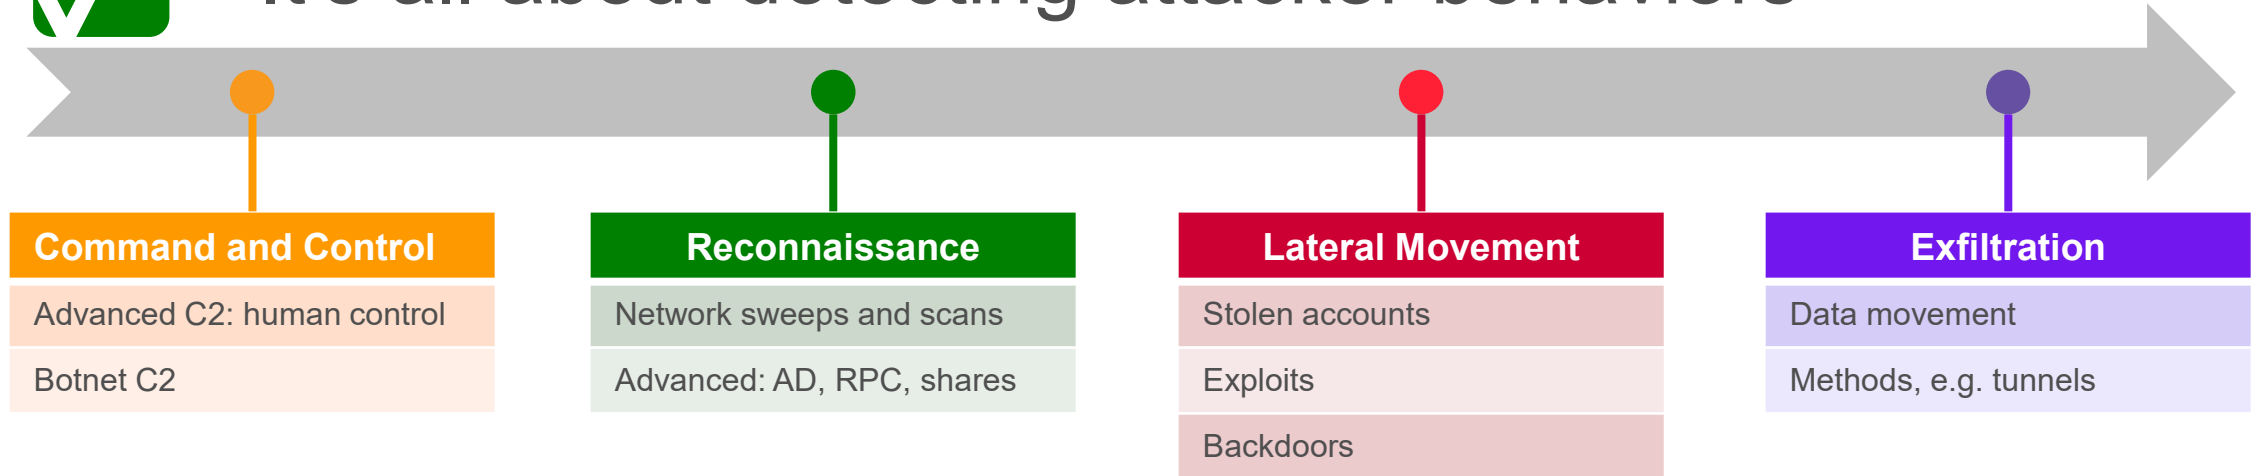


Data Science

- Determine best approach to identify behavior
- Develop and tune models



It's all about detecting attacker behaviors



Security Research

- Identify, prioritize, and characterize fundamental attacker behaviors
- Validate models



Data Science

- Determine best approach to identify behavior
- Develop and tune models



Supervised deep learning to detect remote access

External Remote Access



- Traffic to a multi-dimensional time series
- Deep learning model to featurize the data flow
- Discovers the human on the outside taking control



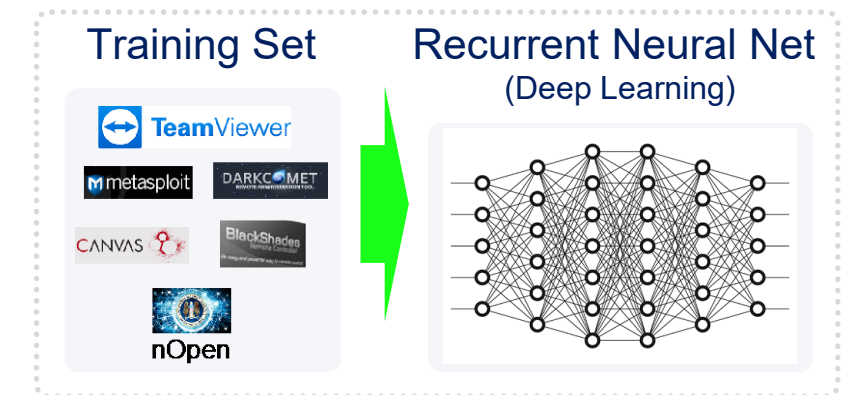
JQSnicker



Security Research



Data Science





Unsupervised learning to detect admin misuse

Suspicious Kerberos Client
Suspicious Admin
Suspicious Remote Exec


Security Research



Authenticate using
a stolen credential



Administer a host
using the stolen
credential



Move laterally using
credential for remote
execution (RPC)



Data Science

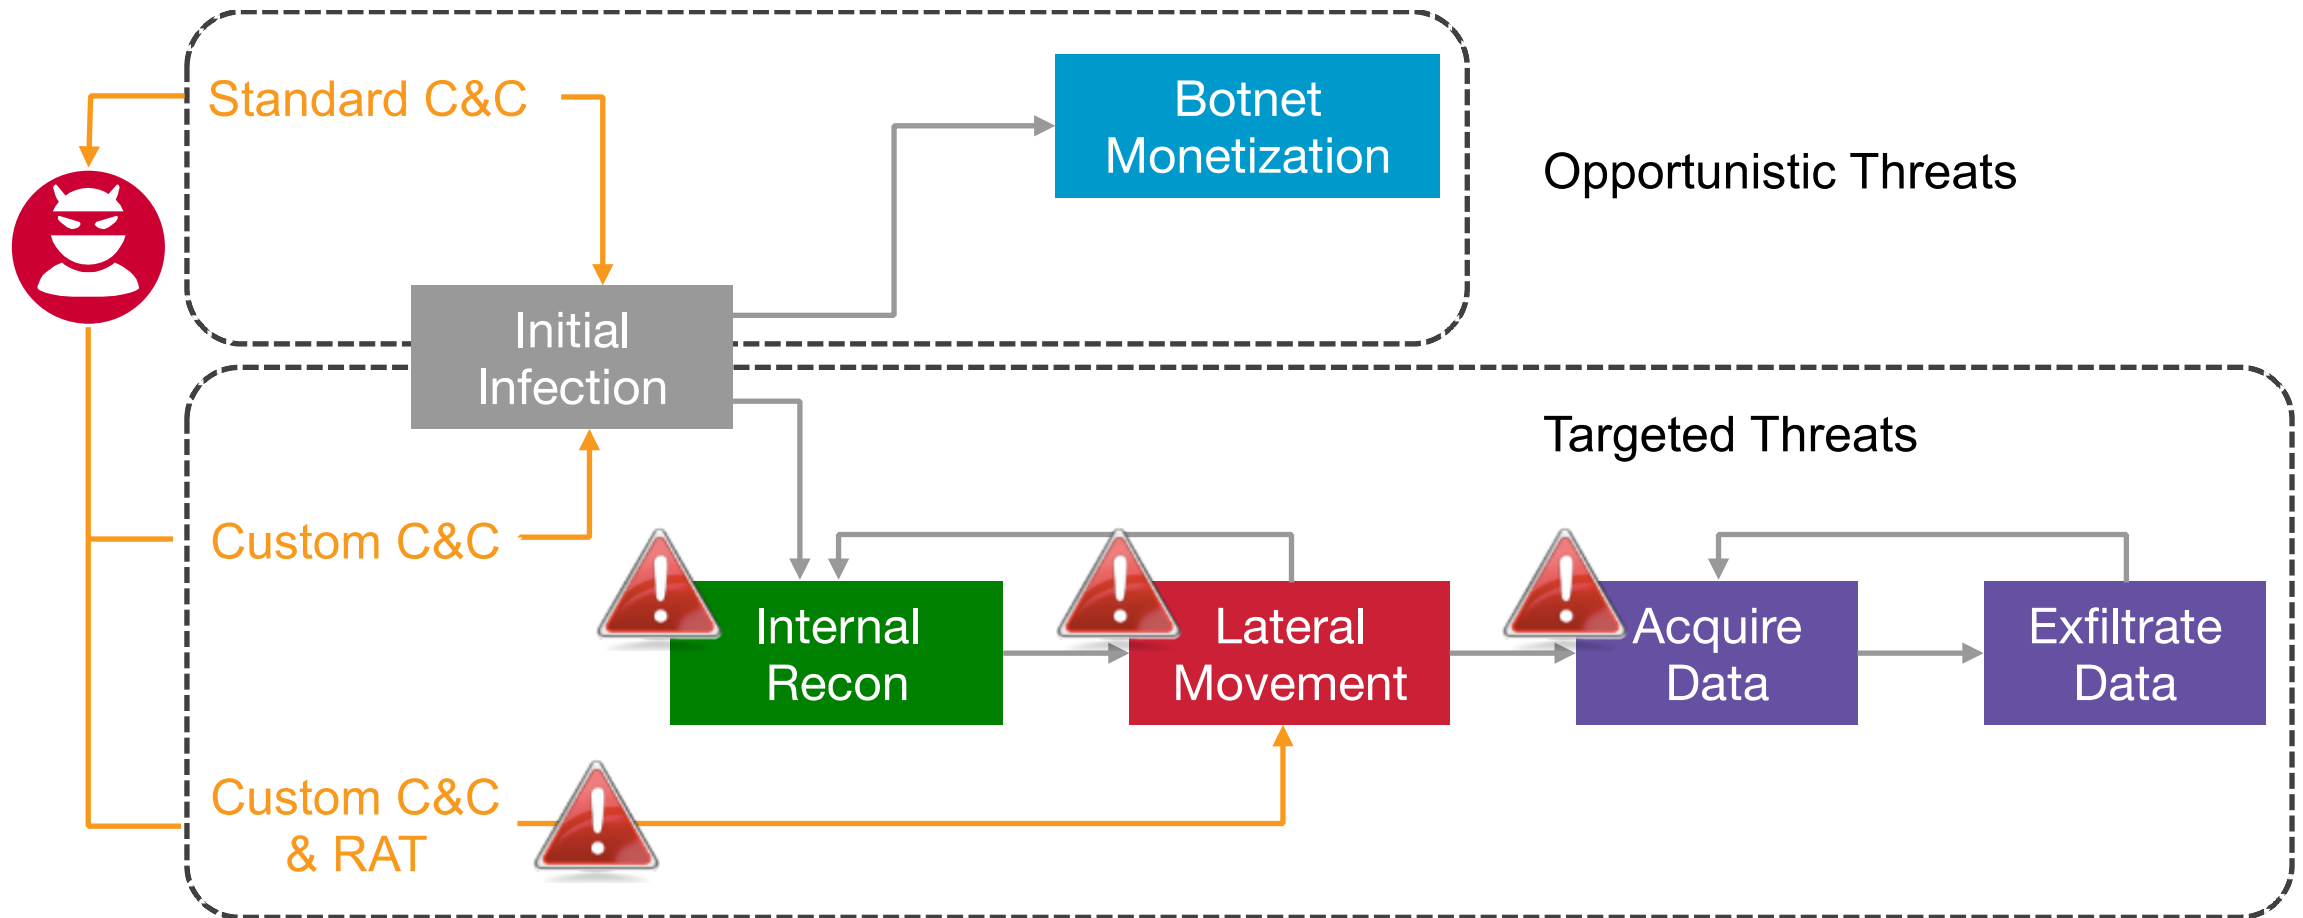
Controller for each host and
identify mismatches

Learn which systems each host
administers, via which protocols, and
identify abnormal administration

Learn normal RPC usage (target, UUID,
named pipe, account tuples) for each
host and identify abnormal usage

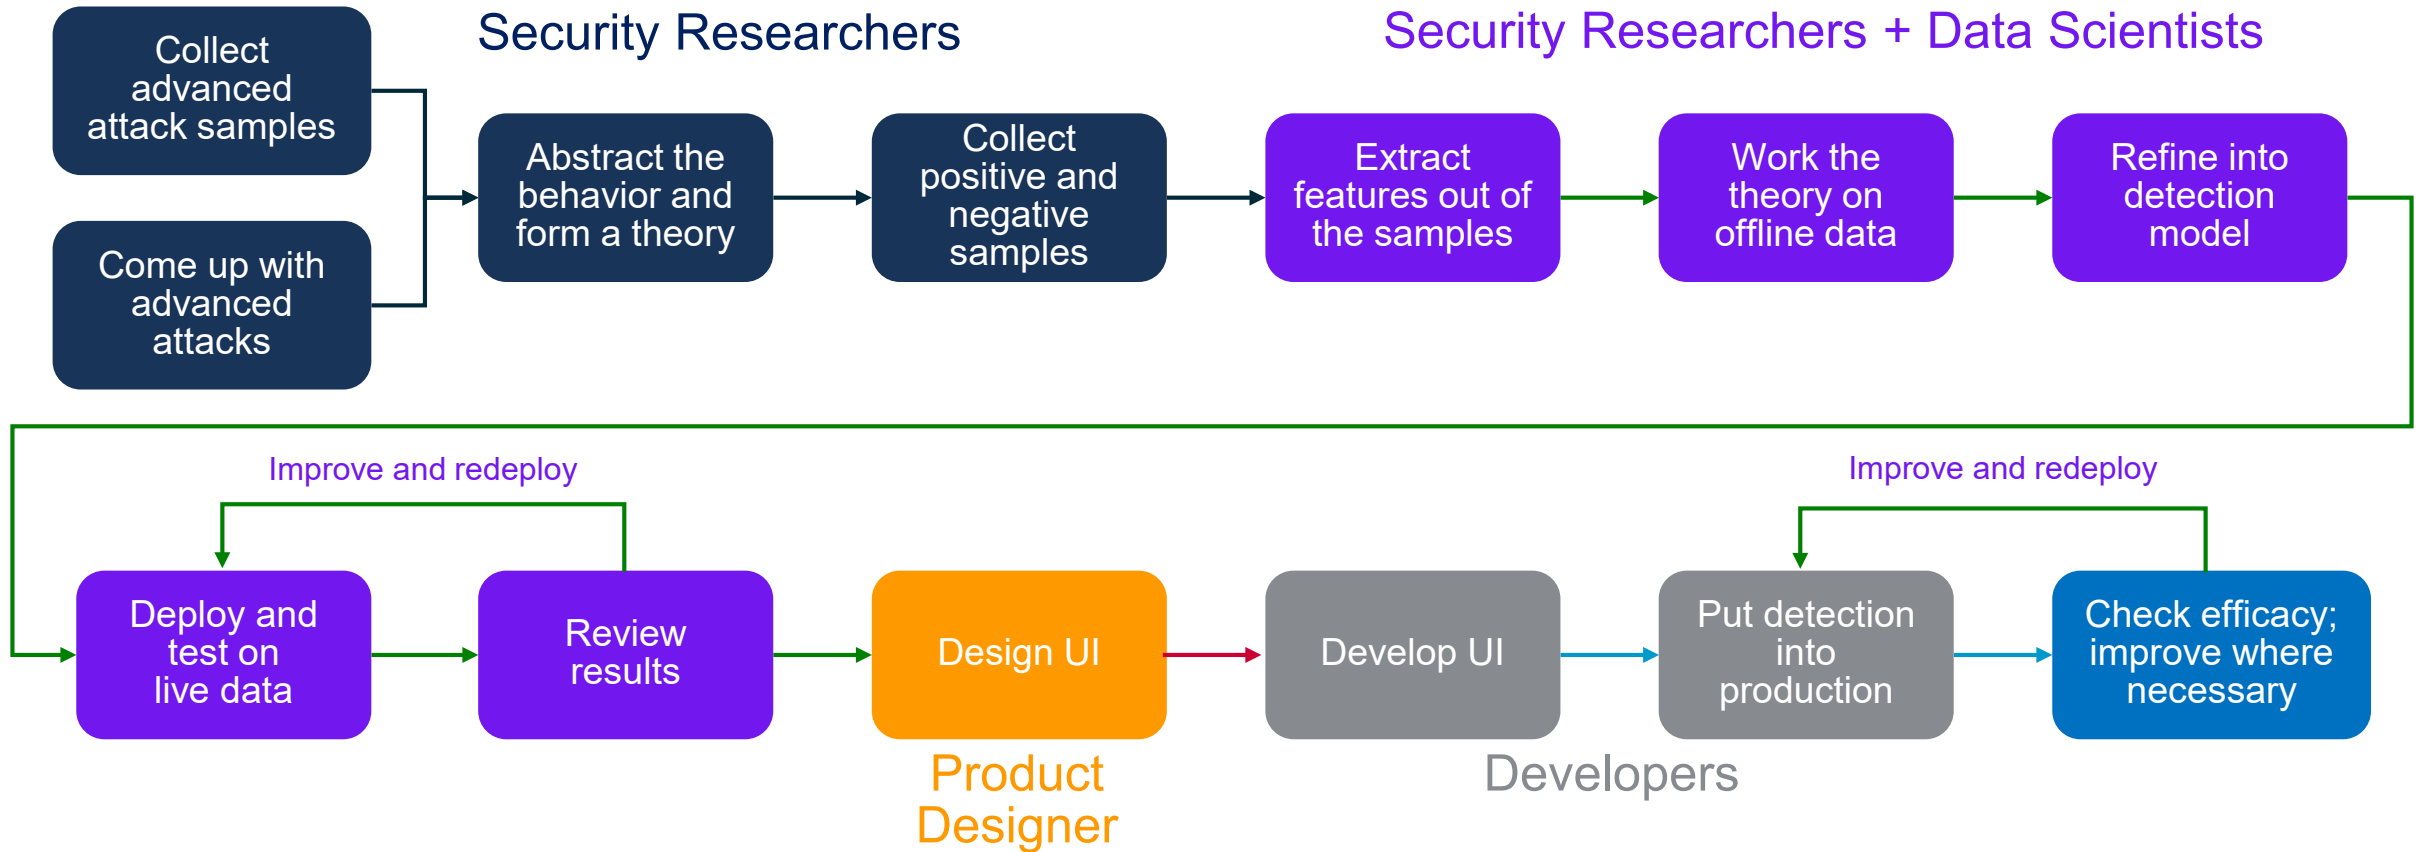


Prioritizing incidents using a time series of events





What it takes to build an algorithm





What to ask when applying AI to threat detection

- What type of machine learning algorithms are used?
- How many machine learning algorithms are applied, and how are they categorized?
- How frequently are algorithms updated and new algorithms released?
- How many algorithms require a learning period, and how long does that take?
- How are critical and high-risk threats prioritized?
- What kind of efficacy can be expected?



Thank you

Chris Morales
chris@Vectra.ai